

# Apache Hadoop – A course for undergraduates

## Lecture 5



# Practical Development Tips and Techniques

---

## Chapter 5.1



# Practical Development Tips and Techniques

---

- **Strategies for debugging MapReduce code**
- **How to test MapReduce code locally using LocalJobRunner**
- **How to write and view log files**
- **How to retrieve job information with counters**
- **Why reusing objects is a best practice**
- **How to create Map-only MapReduce jobs**

# Chapter Topics

---

## Practical Development Tips and Techniques

- **Strategies for Debugging MapReduce Code**
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs

# Introduction to Debugging

---

- **Debugging MapReduce code is difficult!**
  - Each instance of a Mapper runs as a separate task
    - Often on a different machine
  - Difficult to attach a debugger to the process
  - Difficult to catch 'edge cases'
- **Very large volumes of data mean that unexpected input is likely to appear**
  - Code which expects all data to be well-formed is likely to fail

# Common-Sense Debugging Tips

---

- **Code defensively**
  - Ensure that input data is in the expected format
  - Expect things to go wrong
  - Catch exceptions
- **Start small, build incrementally**
- **Make as much of your code as possible Hadoop-agnostic**
  - Makes it easier to test
- **Write unit tests**
- **Test locally whenever possible**
  - With small amounts of data
- **Then test in pseudo-distributed mode**
- **Finally, test on the cluster**

# Testing Strategies

---

- **When testing in pseudo-distributed mode, ensure that you are testing with a similar environment to that on the real cluster**
  - Same amount of RAM allocated to the task JVMs
  - Same version of Hadoop
  - Same version of Java
  - Same versions of third-party libraries



# Chapter Topics

---

## Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- **Testing MapReduce Code Locally Using LocalJobRunner**
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs



## Testing Locally (1)

---

- **Hadoop can run MapReduce in a single, local process**
  - Does not require any Hadoop daemons to be running
  - Uses the local filesystem instead of HDFS
  - Known as LocalJobRunner mode
- **This is a very useful way of quickly testing incremental changes to code**

## Testing Locally (2)

- To run in LocalJobRunner mode, add the following lines to the driver code:

```
Configuration conf = new Configuration();  
conf.set("mapred.job.tracker", "local");  
conf.set("fs.default.name", "file:///");
```

- Or set these options on the command line if your driver uses ToolRunner
  - **fs** is equivalent to **-D fs.default.name**
  - **jt** is equivalent to **-D mapred.job.tracker**
  - e.g.

```
$ hadoop jar myjar.jar MyDriver -fs=file:/// -jt=local \  
indir outdir
```

## Testing Locally (3)

---

- **Some limitations of LocalJobRunner mode:**

- Distributed Cache does not work
- The job can only specify a single Reducer
- Some 'beginner' mistakes may not be caught
  - For example, attempting to share data between Mappers will work, because the code is running in a single JVM

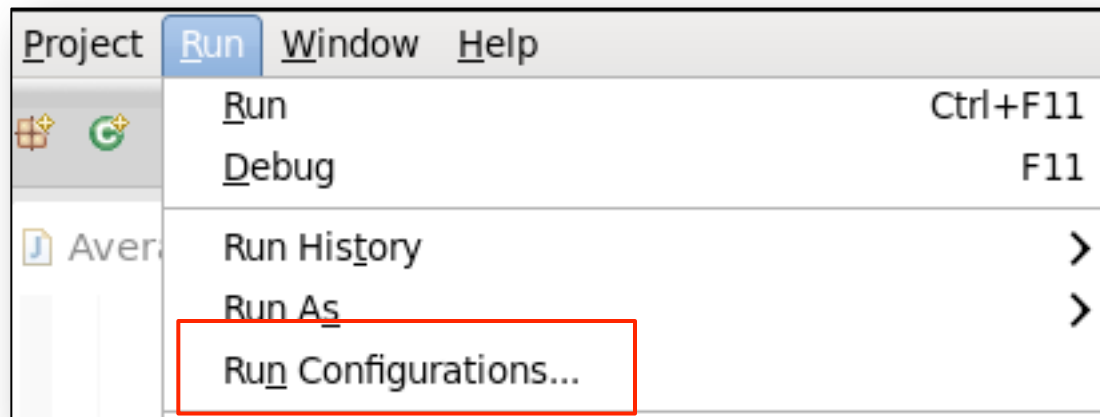
## LocalJobRunner Mode in Eclipse (1)

---

- **Eclipse on the course VM runs Hadoop code in LocalJobRunner mode from within the IDE**
  - This is Hadoop's default behavior when no configuration is provided
- **This allows rapid development iterations**
  - 'Agile programming'

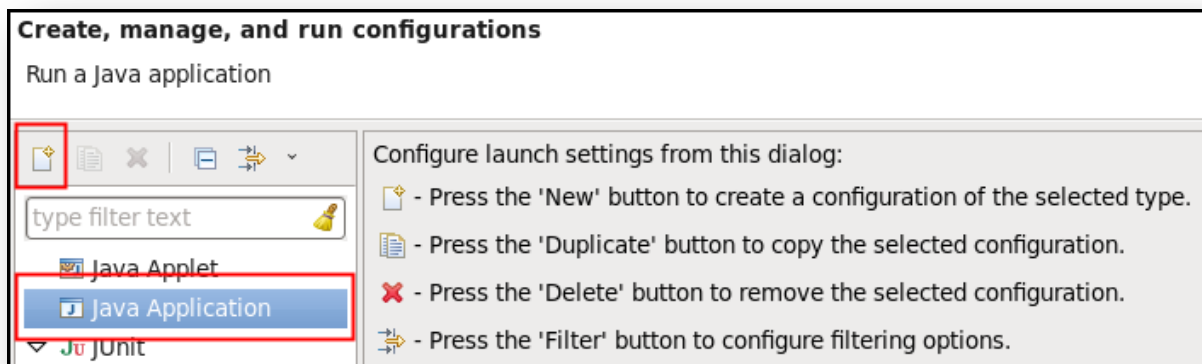
## LocalJobRunner Mode in Eclipse (2)

- Specify a Run Configuration

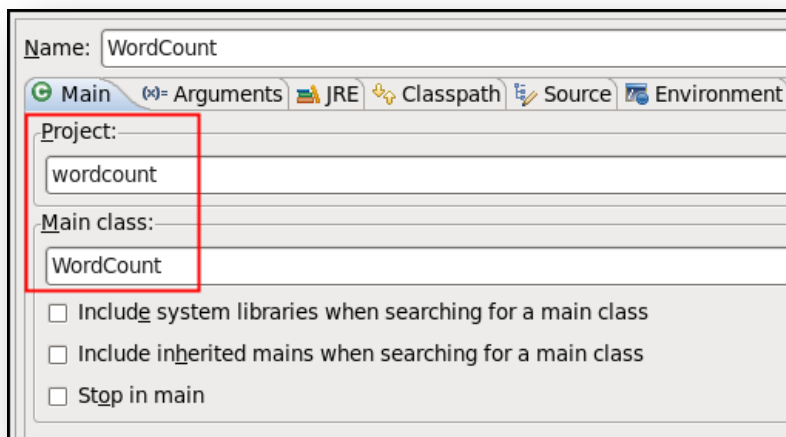


## LocalJobRunner Mode in Eclipse (3)

- **Select Java Application, then select the New button**

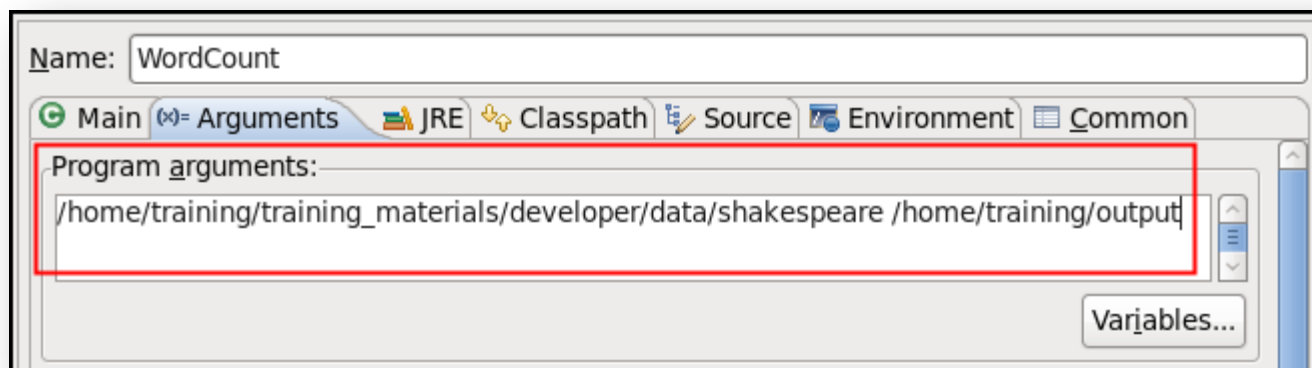


- **Verify that the Project and Main Class fields are pre-filled correctly**




## LocalJobRunner Mode in Eclipse (4)

- **Specify values in the Arguments tab**
  - Local input and output files
  - Any configuration options needed when your job runs



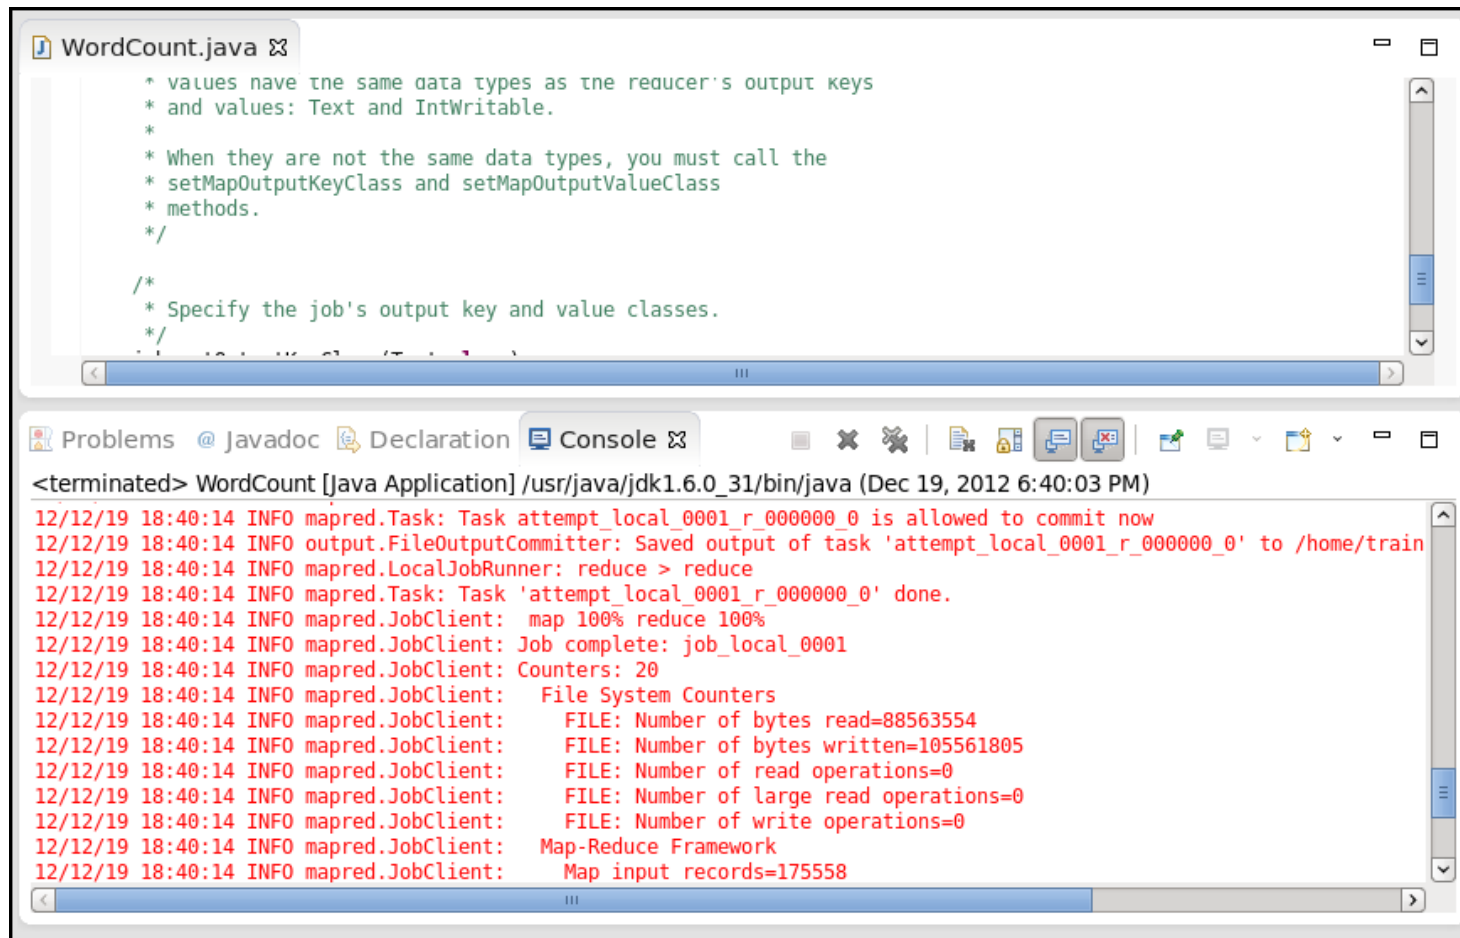
- **Define breakpoints if desired**
- **Execute the application in run mode or debug mode**





## LocalJobRunner Mode in Eclipse (5)

### ■ Review output in the Eclipse console window



The screenshot shows the Eclipse IDE interface. The top editor window displays the `WordCount.java` file with the following comments:

```
/* values have the same data types as the reducer's output keys
 * and values: Text and IntWritable.
 *
 * When they are not the same data types, you must call the
 * setMapOutputKeyClass and setMapOutputValueClass
 * methods.
 */

/* Specify the job's output key and value classes.
 */
```

The bottom console window shows the output of the `WordCount` application. The output is as follows:

```
<terminated> WordCount [Java Application] /usr/java/jdk1.6.0_31/bin/java (Dec 19, 2012 6:40:03 PM)
12/12/19 18:40:14 INFO mapred.Task: Task attempt_local_0001_r_000000_0 is allowed to commit now
12/12/19 18:40:14 INFO output.FileOutputCommitter: Saved output of task 'attempt_local_0001_r_000000_0' to /home/train
12/12/19 18:40:14 INFO mapred.LocalJobRunner: reduce > reduce
12/12/19 18:40:14 INFO mapred.Task: Task 'attempt_local_0001_r_000000_0' done.
12/12/19 18:40:14 INFO mapred.JobClient: map 100% reduce 100%
12/12/19 18:40:14 INFO mapred.JobClient: Job complete: job_local_0001
12/12/19 18:40:14 INFO mapred.JobClient: Counters: 20
12/12/19 18:40:14 INFO mapred.JobClient: File System Counters
12/12/19 18:40:14 INFO mapred.JobClient: FILE: Number of bytes read=88563554
12/12/19 18:40:14 INFO mapred.JobClient: FILE: Number of bytes written=105561805
12/12/19 18:40:14 INFO mapred.JobClient: FILE: Number of read operations=0
12/12/19 18:40:14 INFO mapred.JobClient: FILE: Number of large read operations=0
12/12/19 18:40:14 INFO mapred.JobClient: FILE: Number of write operations=0
12/12/19 18:40:14 INFO mapred.JobClient: Map-Reduce Framework
12/12/19 18:40:14 INFO mapred.JobClient: Map input records=175558
```

# Chapter Topics

---

## Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- **Writing and Viewing Log Files**
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs

## Before Logging: `stdout` and `stderr`

---

- **Tried-and-true debugging technique: write to `stdout` or `stderr`**
- **If running in LocalJobRunner mode, you will see the results of `System.err.println()`**
- **If running on a cluster, that output will not appear on your console**
  - Output is visible via Hadoop's Web UI

## Aside: The Hadoop Web UI

---

- **All Hadoop daemons contain a Web server**
  - Exposes information on a well-known port
- **Most important for developers is the JobTracker Web UI**
  - `http://<job_tracker_address>:50030/`
  - `http://localhost:50030/` if running in pseudo-distributed mode
- **Also useful: the NameNode Web UI**
  - `http://<name_node_address>:50070/`

## Aside: The Hadoop Web UI (cont'd)

- Your instructor will now demonstrate the JobTracker UI

Applications Places System > Mozilla Firefox

File Edit View History Bookmarks Tools Help

http://node1:50030/taskdetails.jsp?jobid=job\_201009070723\_0017&tipid=task\_2

Hadoop Task Details

### Job job\_201009070723\_0017

#### All Task Attempts

Task Attempts	Machine	Status	Progress	Start Time	Finish Time	Errors	Task Logs	Counter	Actions
attempt_201009070723_0017_m_000000_0	/default-rack/node4	SUCCEEDED	100.00%	9-Oct-2010 17:38:40	9-Oct-2010 17:39:02 (21sec)		<a href="#">Last 4KB</a> <a href="#">Last 8KB</a> <a href="#">All</a>	8	

#### Input Split Locations

[Go back to the job](#)  
[Go back to JobTracker](#)

[Cloudera's Distribution for Hadoop, 2010.](#)

Done

ian@node1: ~ Hadoop Task Details - ...

# Logging: Better Than Printing

---

- **`println` statements rapidly become awkward**

- Turning them on and off in your code is tedious, and leads to errors



- **Logging provides much finer-grained control over:**

- What gets logged
  - When something gets logged
  - How something is logged

## Logging With log4j

- Hadoop uses log4j to generate all its log files
- Your Mappers and Reducers can also use log4j
  - All the initialization is handled for you by Hadoop
- Add the log4j.jar-<version> file from your CDH distribution to your classpath when you reference the log4j classes

```
import org.apache.log4j.Level;
import org.apache.log4j.Logger;

class FooMapper implements Mapper {
    private static final Logger LOGGER =
        Logger.getLogger (FooMapper.class.getName());
    ...
}
```



## Logging With `log4j` (cont'd)

- Simply send strings to loggers tagged with severity levels:

```
LOGGER.trace("message");  
LOGGER.debug("message");  
LOGGER.info("message");  
LOGGER.warn("message");  
LOGGER.error("message");
```

- Beware expensive operations like concatenation

- To avoid performance penalty, make it conditional like this:

```
if (LOGGER.isDebugEnabled()) {  
    LOGGER.debug("Account info:" + acct.getReport());  
}
```

# log4j Configuration

- Node-wide configuration for log4j is stored in `/etc/hadoop/conf/log4j.properties`
- Override settings for your application in your own `log4j.properties`
  - Can change global log settings with `hadoop.root.log` property
  - Can override log level on a per-class basis, e.g.

```
log4j.logger.org.apache.hadoop.mapred.JobTracker=WARN
```

```
log4j.logger.com.mycompany.myproject.FooMapper=DEBUG
```

Full class name



- Or set the level programmatically:

```
LOGGER.setLevel(Level.WARN);
```

## Setting Logging Levels for a Job

- You can tell Hadoop to set logging levels for a job using configuration properties

- mapred.map.child.log.level
  - mapred.reduce.child.log.level

- Examples

- Set the logging level to DEBUG for the Mapper

```
$ hadoop jar myjob.jar MyDriver \  
-Dmapred.map.child.log.level=DEBUG indir outdir
```

- Set the logging level to WARN for the Reducer

```
$ hadoop jar myjob.jar MyDriver \  
-Dmapred.reduce.child.log.level=WARN indir outdir
```

# Where Are Log Files Stored?

---

- **Log files are stored on the machine where the task attempt ran**
  - Location is configurable
  - By default:

```
/var/log/hadoop-0.20-mapreduce/  
userlogs/${task.id}/syslog
```
- **You will often not have ssh access to a node to view its logs**
  - Much easier to use the JobTracker Web UI
    - Automatically retrieves and displays the log files for you

# Restricting Log Output

- **If you suspect the input data of being faulty, you may be tempted to log the (key, value) pairs your Mapper receives**
  - Reasonable for small amounts of input data
  - Caution! If your job runs across 500GB of input data, you could be writing up to 500GB of log files!
  - Remember to think at scale...
- **Instead, wrap vulnerable sections of code in `try { . . . }` blocks**
  - Write logs in the `catch { . . . }` block
    - This way only critical data is logged




## Aside: Throwing Exceptions

---

- **You could throw exceptions if a particular condition is met**
  - For example, if illegal data is found

```
throw new RuntimeException("Your message here");
```



- **Usually not a good idea**
  - Exception causes the task to fail
  - If a task fails four times, the entire job will fail

# Chapter Topics

---

## Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- **Retrieving Job Information with Counters**
- Reusing Objects
- Creating Map-only MapReduce Jobs



# What Are Counters? (1)

---

- **Counters provide a way for Mappers or Reducers to pass aggregate values back to the driver after the job has completed**
  - Their values are also visible from the JobTracker's Web UI
  - And are reported on the console when the job ends
- **Very basic: just have a name and a value**
  - Value can be incremented within the code
- **Counters are collected into Groups**
  - Within the group, each Counter has a name
- **Example: A group of Counters called RecordType**
  - Names: `TypeA`, `TypeB`, `TypeC`
  - Appropriate Counter can be incremented as each record is read in the Mapper

## What Are Counters? (2)

---

- Counters can be set and incremented via the method

```
context.getCounter(group, name).increment(amount);
```

- Example:

```
context.getCounter("RecordType", "A").increment(1);
```

## Retrieving Counters in the Driver Code

- To retrieve Counters in the Driver code after the job is complete, use code like this in the driver:

```
long typeARecords =  
    job.getCounters().findCounter("RecordType", "A").getValue();  
  
long typeBRecords =  
    job.getCounters().findCounter("RecordType", "B").getValue();
```

## Counters: Caution

---

- **Do not rely on a counter's value from the Web UI while a job is running**
  - Due to possible speculative execution, a counter's value could appear larger than the actual final value
  - Modifications to counters from subsequently killed/failed tasks will be removed from the final count

# Chapter Topics

---

## Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- **Reusing Objects**
- Creating Map-only MapReduce Jobs

## Reuse of Objects is Good Practice (1)

- It is generally good practice to reuse objects
  - Instead of creating many new objects
- Example: Our original WordCount Mapper code

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String word = value.toString();
        for (int i = 0; i < word.length(); i++) {
            if (i > 0) {
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Each time the `map()` method is called, we create a new `Text` object and a new `IntWritable` object.

## Reuse of Objects is Good Practice (2)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException
    {
        String line = value.toString();

        for (String word : line.split("\\W+")) {
            if (word.length() > 0) {
                wordObject.set(word);
                context.write(wordObject, one);
            }
        }
    }
}
```

Create objects for the key and value outside of your `map()` method



## Reuse of Objects is Good Practice (3)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();
```

Within the `map()` method, populate the objects and write them out. Hadoop will take care of serializing the data so it is perfectly safe to re-use the objects.

```
    if (word.length() > 0) {
        wordObject.set(word);
        context.write(wordObject, one);
```

```
    }
```

```
}
```

```
}
```

```
}
```

# Object Reuse: Caution!

---

- **Hadoop re-uses objects all the time**
- **For example, each time the Reducer is passed a new value, the same object is reused**
- **This can cause subtle bugs in your code**
  - For example, if you build a list of value objects in the Reducer, each element of the list will point to the same underlying object
    - Unless you do a deep copy

# Chapter Topics

---

## Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- **Creating Map-only MapReduce jobs**

# Map-Only MapReduce Jobs

---

- **There are many types of job where only a Mapper is needed**
- **Examples:**
  - Image processing
  - File format conversion
  - Input data sampling
  - ETL

## Creating Map-Only Jobs

- To create a Map-only job, set the number of Reducers to 0 in your Driver code

```
job.setNumReduceTasks(0);
```

- Call the `Job.setOutputKeyClass` and `Job.setOutputValueClass` methods to specify the output types
  - *Not* the `Job.setMapOutputKeyClass` and `Job.setMapOutputValueClass` methods
- Anything written using the `Context.write` method in the Mapper will be written to HDFS
  - Rather than written as intermediate data
  - One file per Mapper will be written

## Key Points

---

- **LocalJobRunner lets you test jobs on your local machine**
- **Hadoop uses the Log4J framework for logging**
- **Reusing objects is a best practice**
- **Counters provide a way of passing numeric data back to the driver**
- **Create Map-only MapReduce jobs by setting the number of Reducers to zero**

# Bibliography

---

The following offer more information on topics discussed in this chapter

- **Java version selection**

- <http://wiki.apache.org/hadoop/HadoopJavaVersions>

- **For an example of image processing and file format conversion in Hadoop, see**

- <http://open.blogs.nytimes.com/2007/11/01/self-service-prorated-super-computing-fun/>