



Spark RDD Persistence

Chapter 15



Course Chapters

1	Introduction	Course Introduction
2	Introduction to Hadoop and the Hadoop Ecosystem	Introduction to Hadoop
3	Hadoop Architecture and HDFS	
4	Importing Relational Data with Apache Sqoop	
5	Introduction to Impala and Hive	Importing and Modeling Structured Data
6	Modeling and Managing Data with Impala and Hive	
7	Data Formats	
8	Data File Partitioning	
9	Capturing Data with Apache Flume	Ingesting Streaming Data
10	Spark Basics	Distributed Data Processing with Spark
11	Working with RDDs in Spark	
12	Aggregating Data with Pair RDDs	
13	Writing and Deploying Spark Applications	
14	Parallel Processing in Spark	
15	Spark RDD Persistence	
16	Common Patterns in Spark Data Processing	
17	Spark SQL and DataFrames	
18	Conclusion	Course Conclusion

Spark RDD Persistence

In this chapter you will learn

- **How Spark uses an RDD's lineage in operations**
- **How to persist RDDs to improve performance**

Chapter Topics

Spark RDD Persistence

Distributed Data Processing with Spark

- **RDD Lineage**
- RDD Persistence Overview
- Distributed Persistence
- Conclusion
- Homework: Persist an RDD

Lineage Example (1)

- Each *transformation* operation creates a new *child* RDD



File: purplecow.txt

```
I've never seen a purple cow.  
I never hope to see one;  
But I can tell you, anyhow,  
I'd rather see than be one.
```

Lineage Example (2)

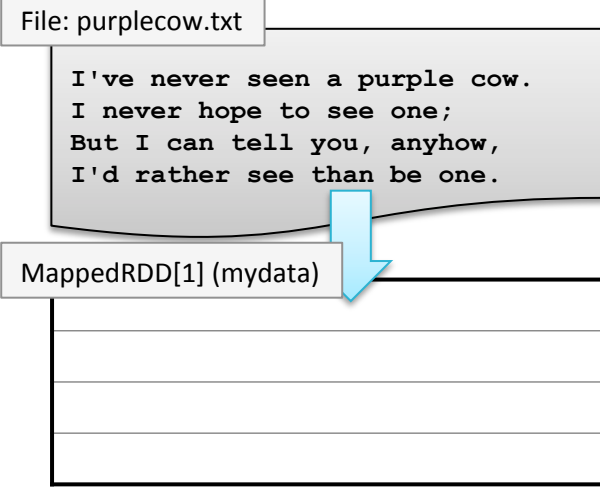
- Each *transformation* operation creates a new *child* RDD

```
> mydata = sc.textFile("purplecow.txt")
```

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

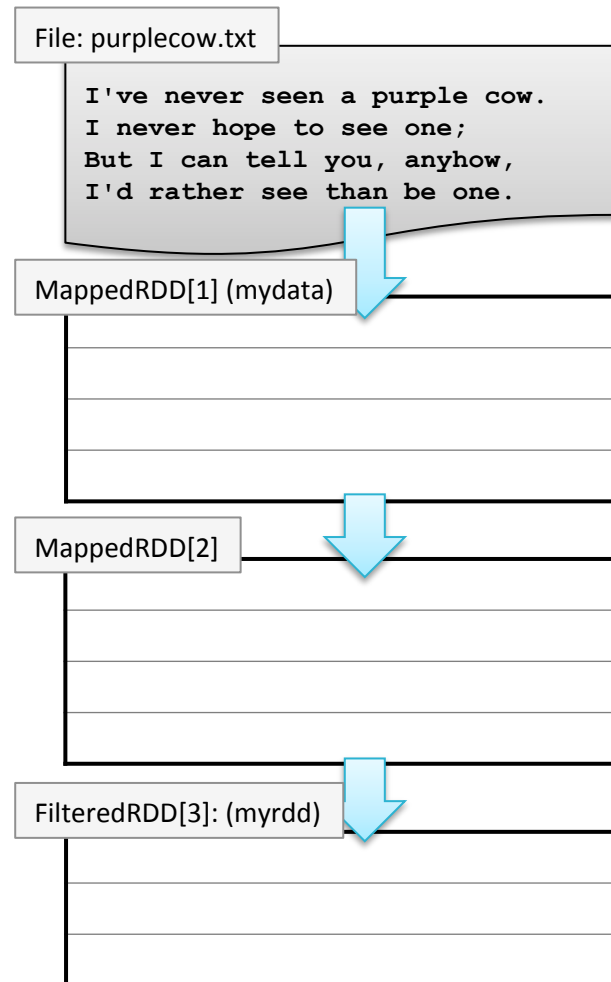
MappedRDD[1] (mydata)



Lineage Example (3)

- Each *transformation* operation creates a new *child* RDD

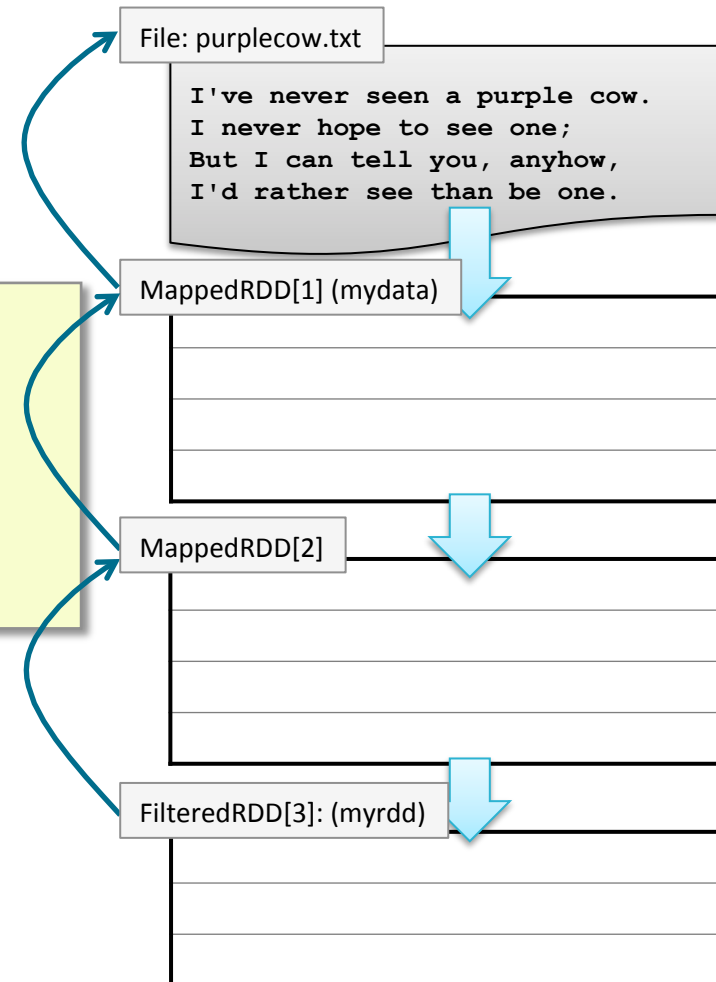
```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s:s.startswith('I'))
```



Lineage Example (4)

- Spark keeps track of the *parent* RDD for each new RDD
- Child RDDs *depend on* their parents

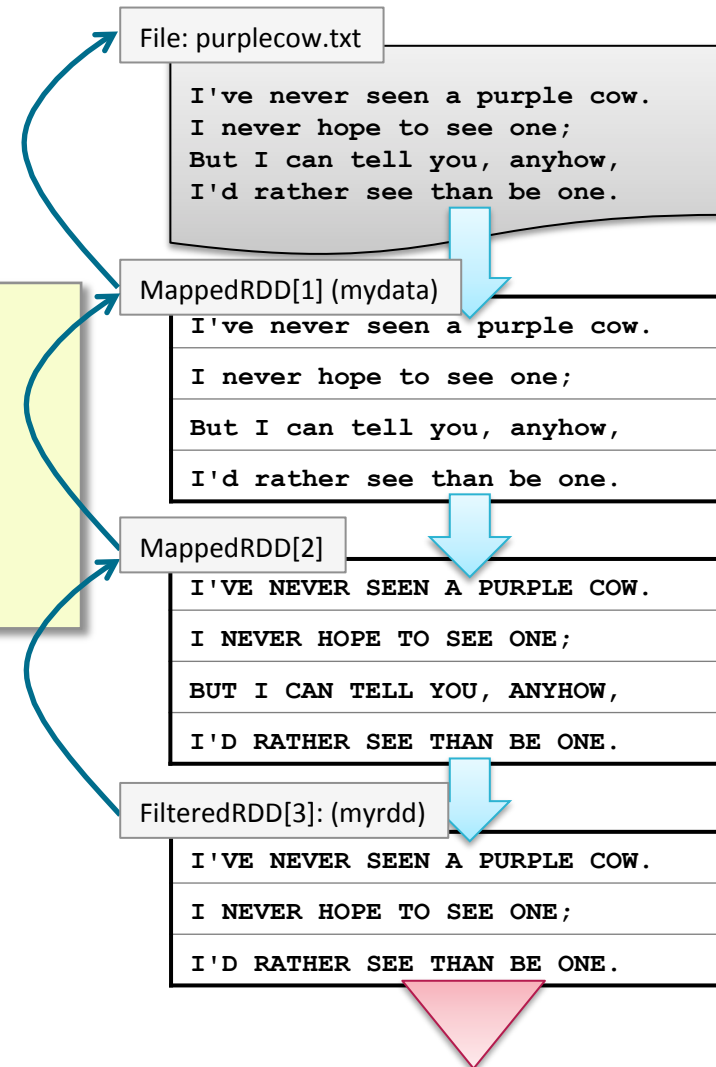
```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper())\
    .filter(lambda s:s.startswith('I'))
```



Lineage Example (5)

- **Action** operations execute the parent transformations

```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper())\
  .filter(lambda s:s.startswith('I'))
> myrdd.count()
3
```

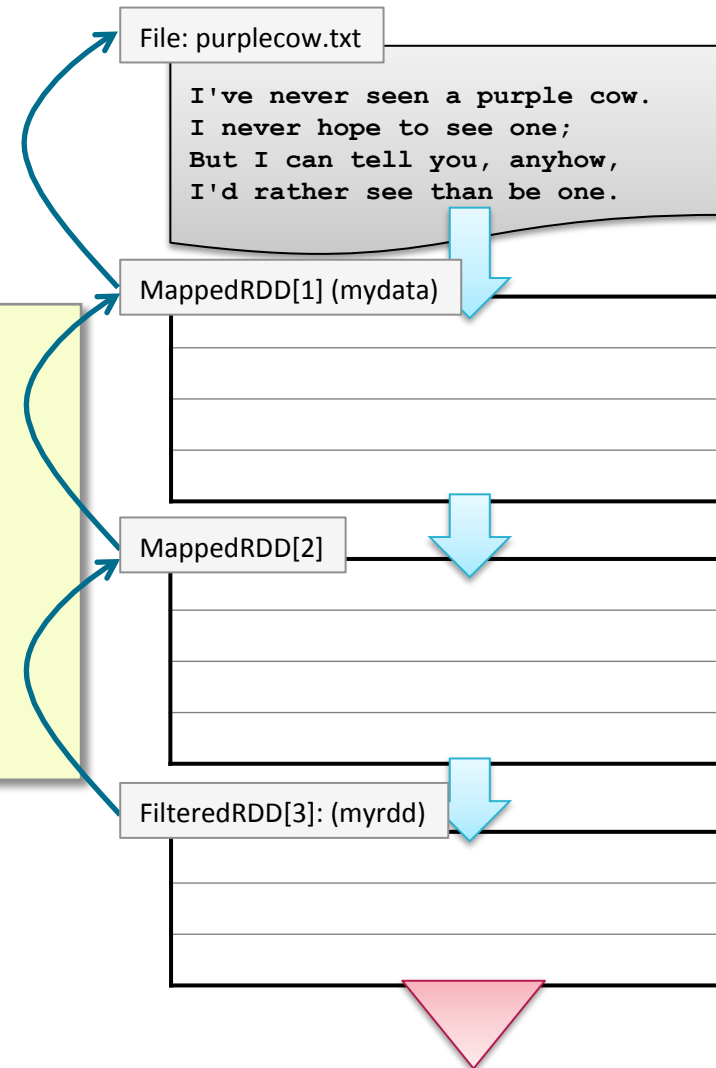


Lineage Example (6)

- Each action re-executes the lineage transformations starting with the base

– By default

```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s:s.startswith('I'))
> myrdd.count()
3
> myrdd.count()
```

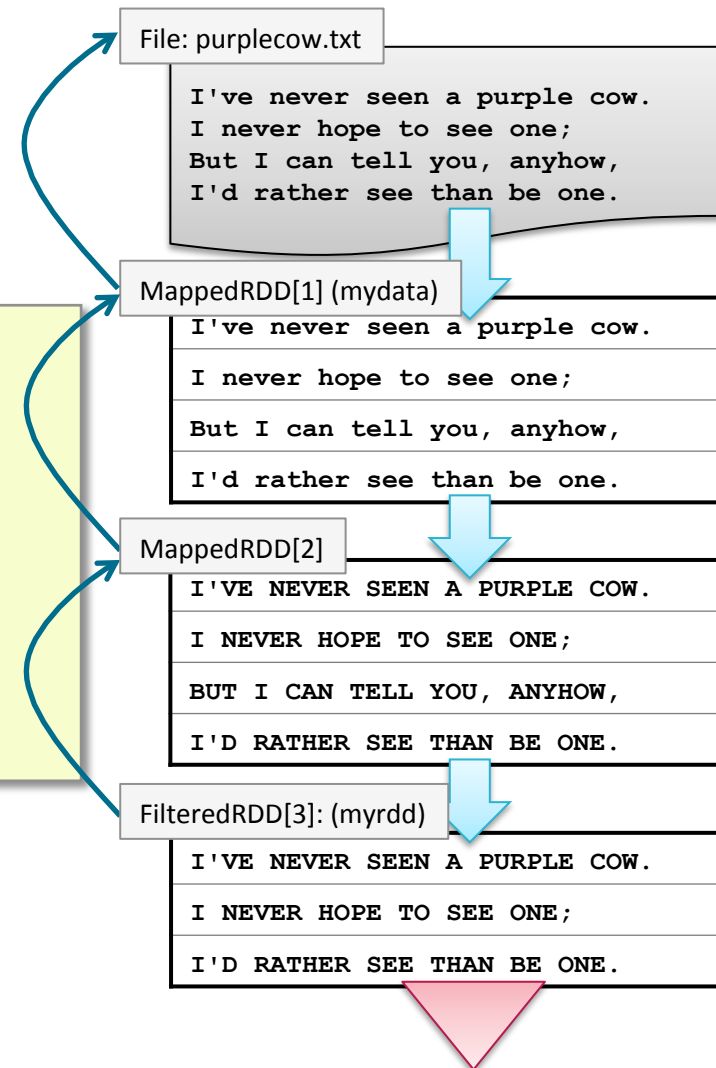


Lineage Example (7)

- Each action re-executes the lineage transformations starting with the base

– By default

```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s: s.startswith('I'))
> myrdd.count()
3
> myrdd.count()
3
```



Chapter Topics

Spark RDD Persistence

Distributed Data Processing with Spark

- RDD Lineage
- **RDD Persistence Overview**
- Distributed Persistence
- Conclusion
- Homework: Persist an RDD

RDD Persistence

- **Persisting an RDD saves the data (by default in memory)**

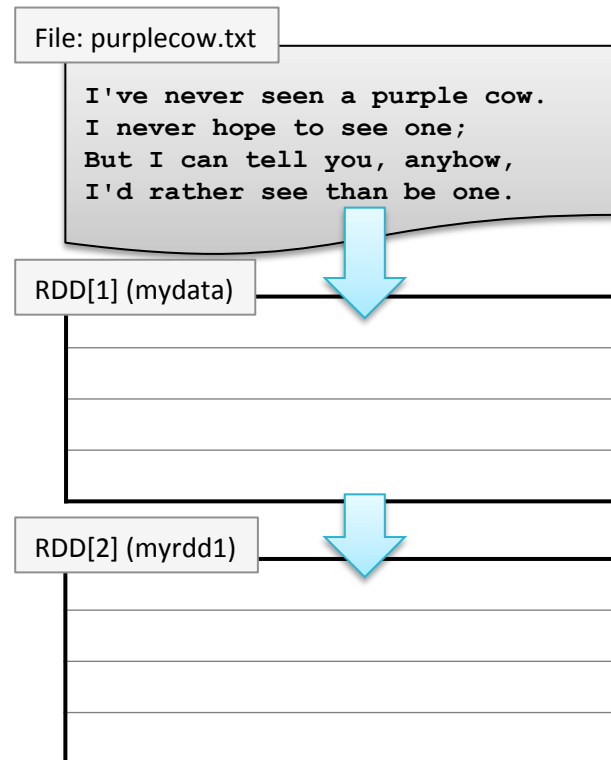
File: purplecow.txt

```
I've never seen a purple cow.  
I never hope to see one;  
But I can tell you, anyhow,  
I'd rather see than be one.
```

RDD Persistence

- Persisting an RDD saves the data (by default in memory)

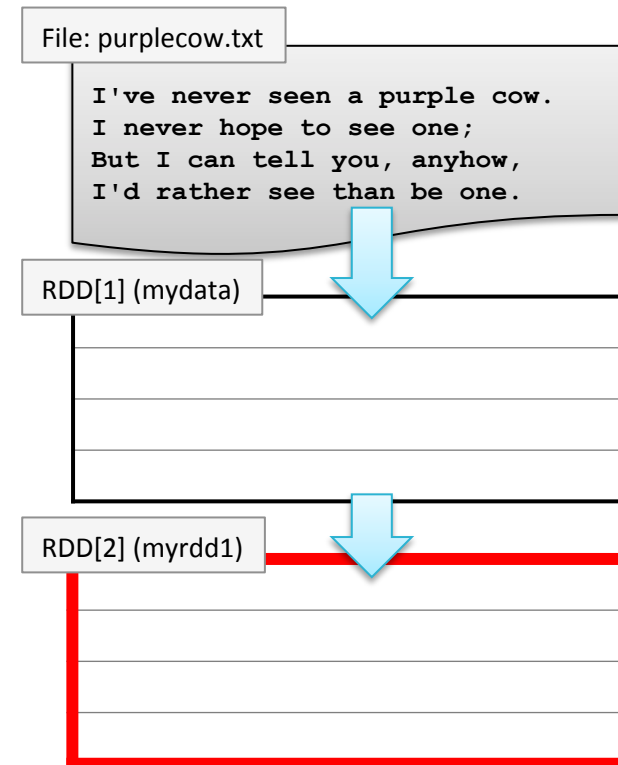
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
```



RDD Persistence

- Persisting an RDD saves the data (by default in memory)

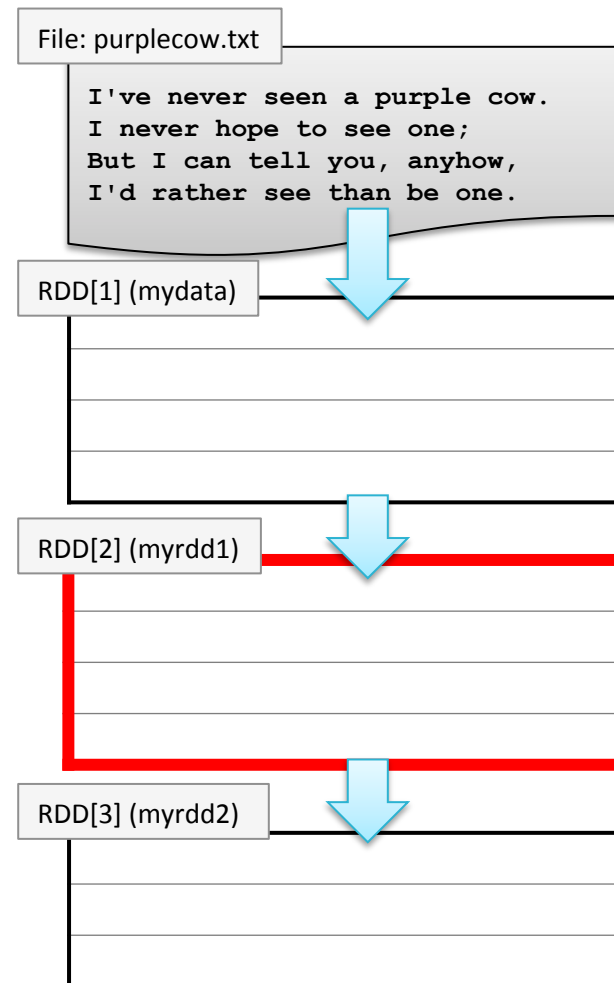
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
> myrdd1.persist()
```



RDD Persistence

- Persisting an RDD saves the data (by default in memory)

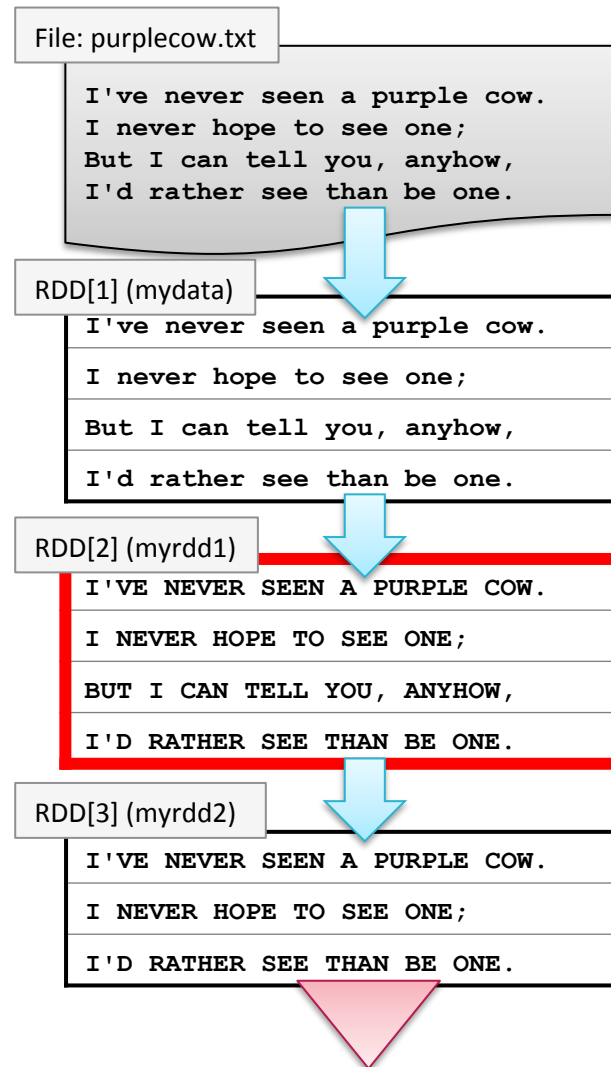
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
> myrdd1.persist()
> myrdd2 = myrdd1.filter(lambda \
    s:s.startswith('I'))
```



RDD Persistence

- Persisting an RDD saves the data (by default in memory)

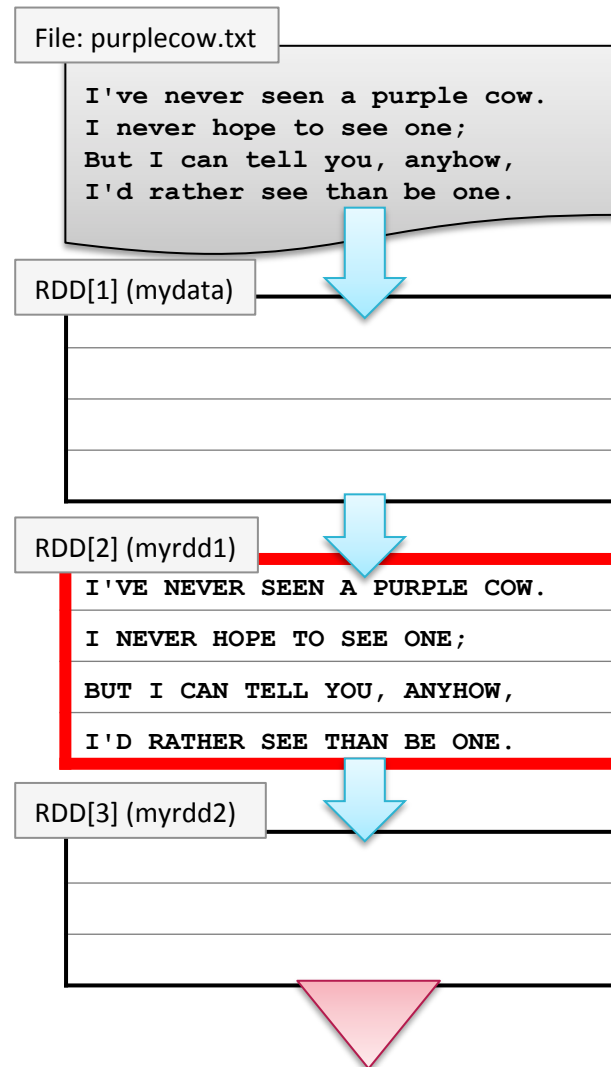
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
> myrdd1.persist()
> myrdd2 = myrdd1.filter(lambda \
    s:s.startswith('I'))
> myrdd2.count()
3
```



RDD Persistence

- Subsequent operations use saved data

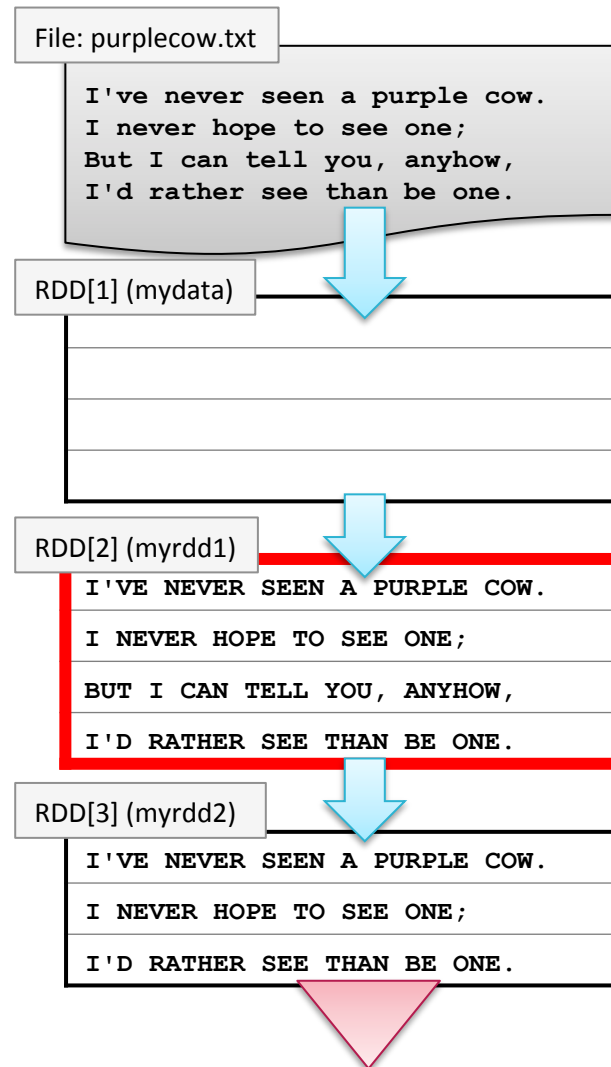
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
> myrdd1.persist()
> myrdd2 = myrdd1.filter(lambda \
    s:s.startswith('I'))
> myrdd2.count()
3
> myrdd2.count()
```



RDD Persistence

- Subsequent operations use saved data

```
> my data =  
  sc.textFile("purplecow.txt")  
> myrdd1 = mydata.map(lambda s:  
  s.upper())  
> myrdd1.persist()  
> myrdd2 = myrdd1.filter(lambda \  
  s:s.startswith('I'))  
> myrdd2.count()  
3  
> myrdd2.count()  
3
```



Memory Persistence

- **In-memory persistence is a *suggestion* to Spark**
 - If not enough memory is available, persisted partitions will be cleared from memory
 - Least recently used partitions cleared first
 - Transformations will be re-executed using the lineage when needed

Chapter Topics

Spark RDD Persistence

Distributed Data Processing with Spark

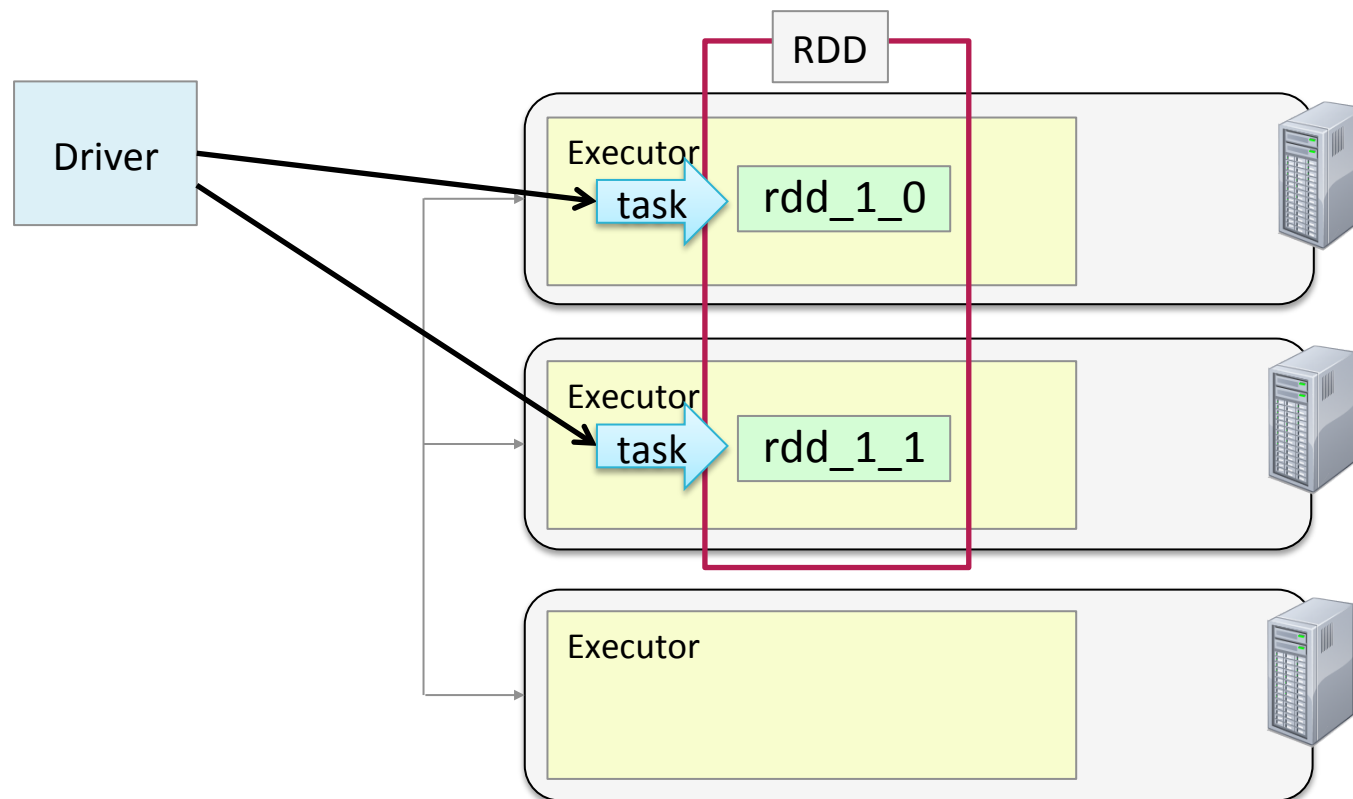
- RDD Lineage
- RDD Persistence Overview
- **Distributed Persistence**
- Conclusion
- Homework: Persist an RDD

Persistence and Fault-Tolerance

- **RDD = *Resilient* Distributed Dataset**
 - Resiliency is a product of tracking lineage
 - RDDs can always be recomputed from their base if needed

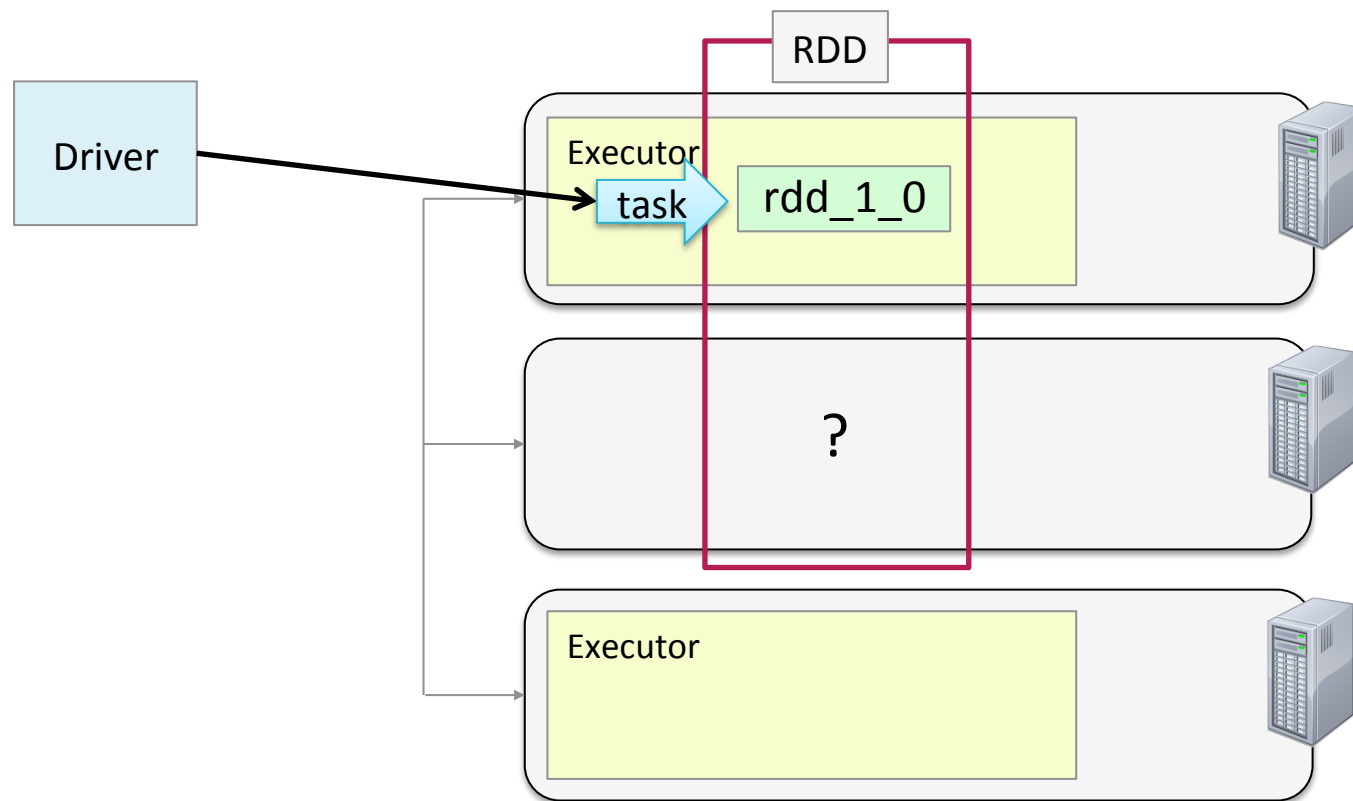
Distributed Persistence

- RDD partitions are distributed across a cluster
- By default, partitions are persisted in memory in Executor JVMs



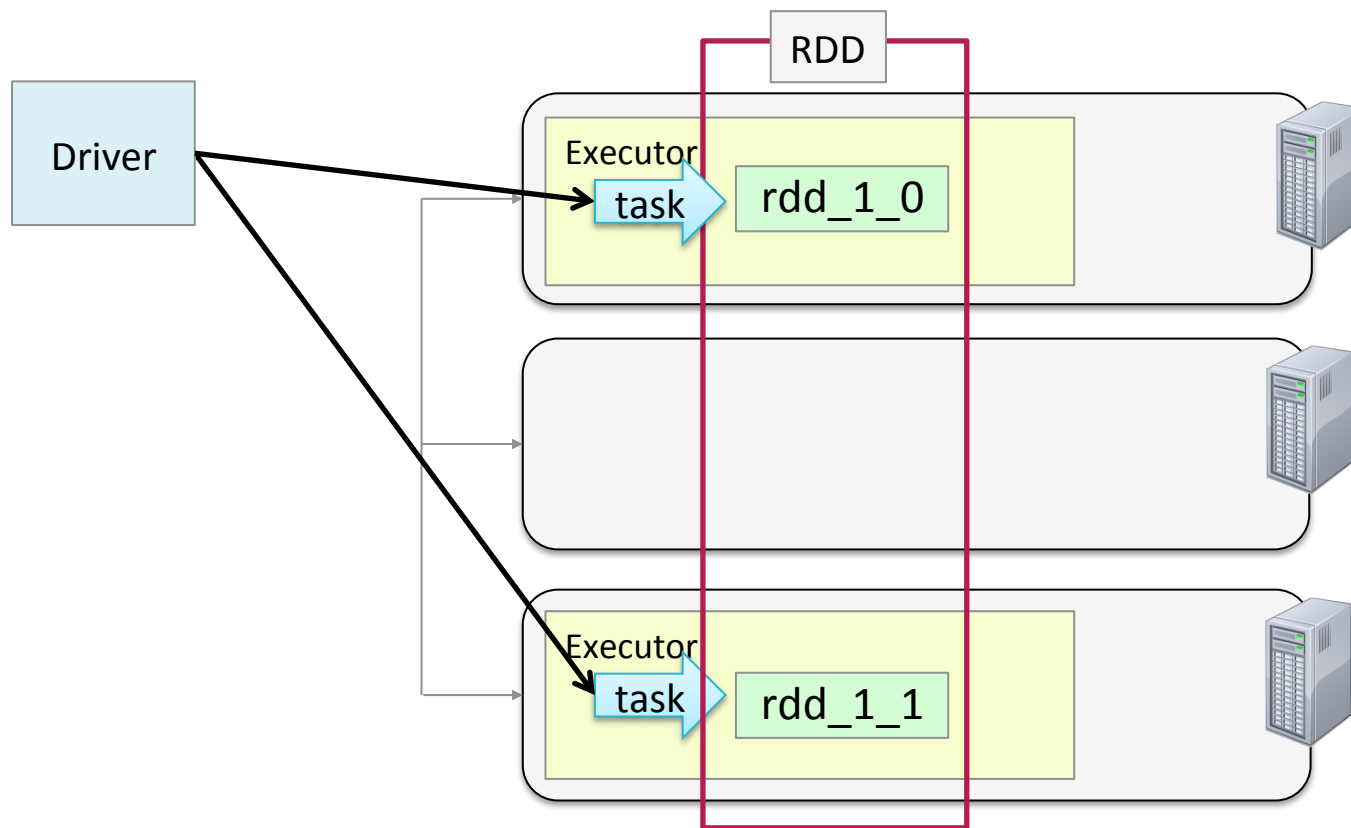
RDD Fault-Tolerance (1)

- What happens if a partition persisted in memory becomes unavailable?



RDD Fault-Tolerance (2)

- The driver starts a new task to recompute the partition on a different node
- Lineage is preserved, data is never lost



Persistence Levels

- **By default, the `persist` method stores data in memory only**
 - The **`cache`** method is a synonym for default (memory) `persist`
- **The `persist` method offers other options called Storage Levels**
- **Storage Levels let you control**
 - Storage location
 - Format in memory
 - Partition replication

Persistence Levels: Storage Location

- **Storage location – where is the data stored?**
 - **MEMORY_ONLY** (default) – same as **cache**
 - **MEMORY_AND_DISK** – Store partitions on disk if they do not fit in memory
 - Called *spilling*
 - **DISK_ONLY** – Store all partitions on disk

Python

```
> from pyspark import StorageLevel  
> myrdd.persist(StorageLevel.DISK_ONLY)
```

Scala

```
> import org.apache.spark.storage.StorageLevel  
> myrdd.persist(StorageLevel.DISK_ONLY)
```

Persistence Levels: Memory Format

- **Serialization – you can choose to serialize the data in memory**
 - **MEMORY_ONLY_SER** and **MEMORY_AND_DISK_SER**
 - Much more space efficient
 - Less time efficient
 - If using Java or Scala, choose a fast serialization library (e.g. Kryo)

Persistence Levels: Partition Replication

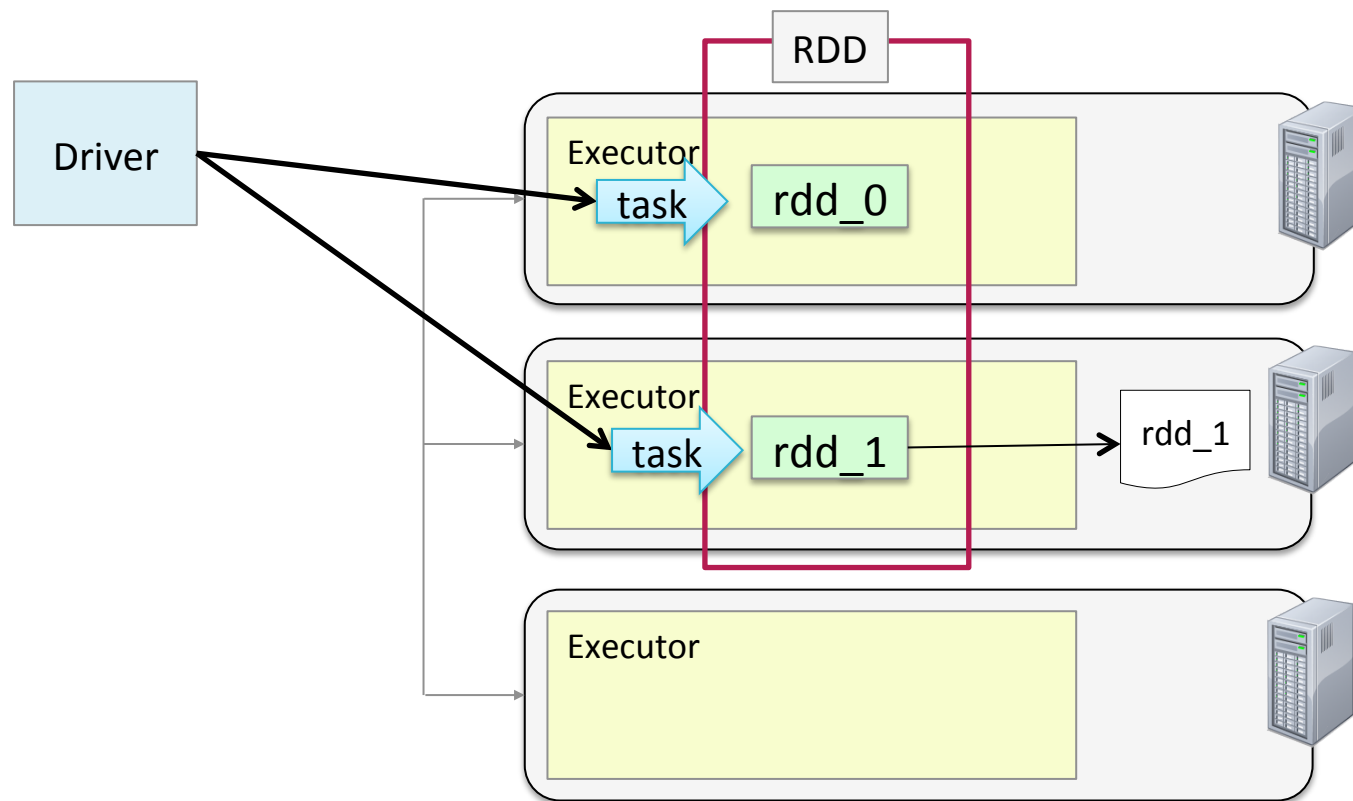
- Replication – store partitions on two nodes
 - MEMORY_ONLY_2
 - MEMORY_AND_DISK_2
 - DISK_ONLY_2
 - MEMORY_AND_DISK_SER_2
 - DISK_ONLY_2
 - You can also define custom storage levels

Changing Persistence Options

- To stop persisting and remove from memory and disk
 - `rdd.unpersist()`
- To change an RDD to a different persistence level
 - Unpersist first

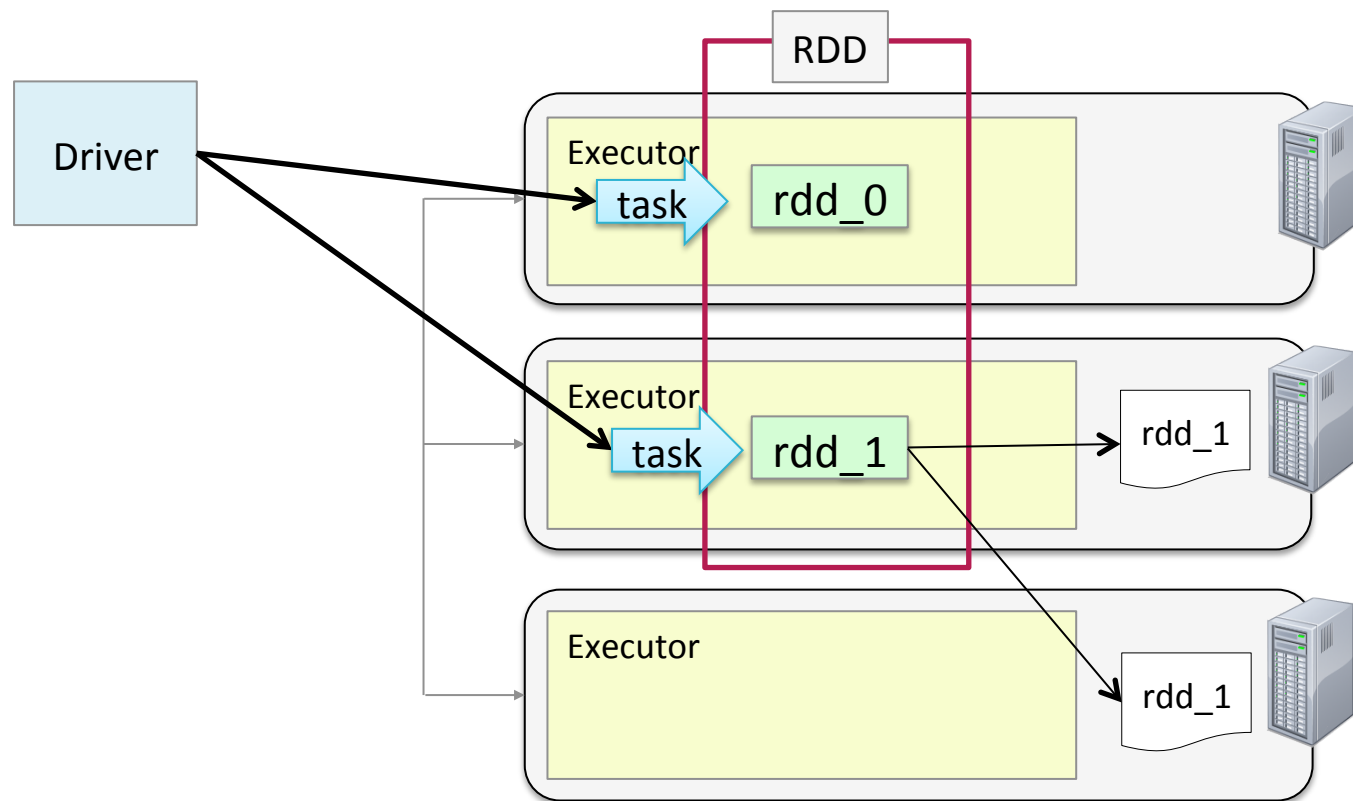
Disk Persistence

- Disk-persisted partitions are stored in local files



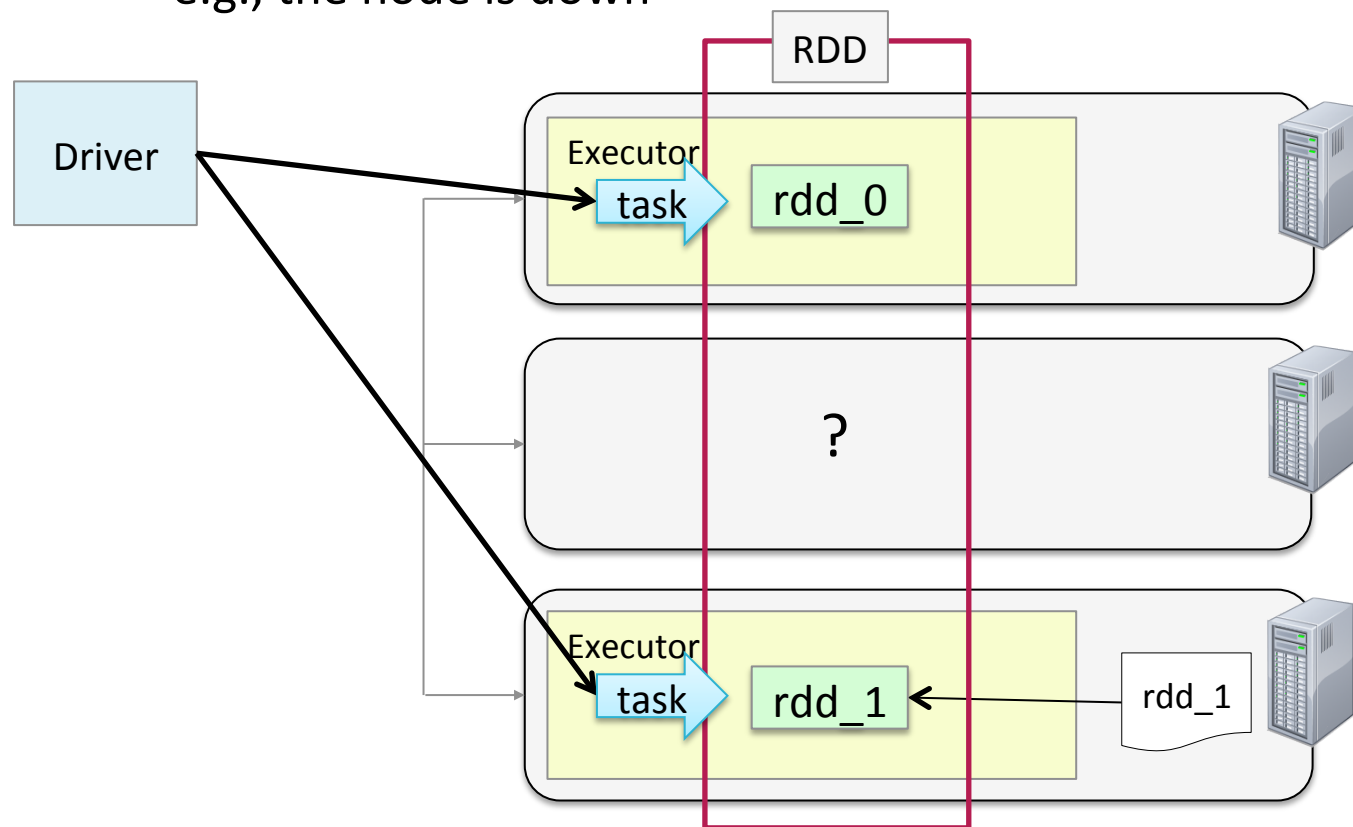
Disk Persistence with Replication (1)

- Persistence replication makes recomputation less likely to be necessary



Disk Persistence with Replication (2)

- **Replicated data on disk will be used to recreate the partition if possible**
 - Will be recomputed if the data is unavailable
 - e.g., the node is down



When and Where to Persist

- **When should you persist a dataset?**

- When a dataset is likely to be re-used
 - e.g., iterative algorithms, machine learning

- **How to choose a persistence level**

- Memory only – when possible, best performance
 - Save space by saving as serialized objects in memory if necessary
- Disk – choose when recomputation is more expensive than disk read
 - e.g., expensive functions or filtering large datasets
- Replication – choose when recomputation is more expensive than memory

Chapter Topics

Spark RDD Persistence

Distributed Data Processing with Spark

- RDD Lineage
- RDD Persistence Overview
- Distributed Persistence
- **Conclusion**
- Homework: Persist an RDD

Essential Points

- **Spark keeps track of each RDD's lineage**
 - Provides fault tolerance
- **By default, every RDD operation executes the entire lineage**
- **If an RDD will be used multiple times, persist it to avoid re-computation**
- **Persistence options**
 - Location – memory only, memory and disk , disk only
 - Format – in-memory data can be serialized to save memory (but at the cost of performance)
 - Replication – saves data on multiple nodes in case a node goes down, for job recovery without recomputation

Chapter Topics

Spark RDD Persistence

Distributed Data Processing with Spark

- RDD Lineage
- RDD Persistence Overview
- Distributed Persistence
- Conclusion
- **Homework: Persist an RDD**

Homework: Persist an RDD

- **In this homework assignment you will**
 - Persist an RDD before reusing it
 - Use the Spark Application UI to see how an RDD is persisted
- **Please refer to the Homework description**